

CPE 431/531

**Chapter 6 – Parallel Processors
from Client to Cloud**

Dr. Rhonda Kay Gaede



6.1 Motivation

- Why multiprocessors?

- _____
- _____
- _____

6.2 The Difficulty of Creating Parallel Programs

- The difficulty with parallelism is not the _____, it's the _____.
- Why is it difficult to write parallel processing programs that are fast?
 - _____
 - _____
 - _____
 - _____
 - _____

6.2 Speedup Challenge

- Suppose you want to achieve a speedup of 90 times faster with 100 processors. What percentage of the original computation can be sequential?

6.2 Speedup Challenge – Bigger Problem

- Suppose you want to perform two sums: one is a sum of 10 scalar variables and one is a matrix sum of a pair of two-dimensional arrays, size 10 by 10. What speedup do you get with 10 versus 40 processors?

6.2 Speedup Challenge – Bigger Problem

- Next, calculate the speed-ups assuming the matrices grow to 20 by 20
- Strong scaling
- Weak scaling

6.2 Speedup Challenge: Balancing Load

- To achieve the speed-up of 20.5 on the previous larger problem with 40 processors, we assumed the load was perfectly balanced (each processor did 2.5 % of the work). Instead, show the impact on speed-up if one processor's load is higher than all the rest. Calculate at 5% and 12.5%.

6.3 SISD, MIMD, SIMD, SPMD, and Vector

- SISD is the normal case – single instruction, single data.
- MIMD – multiple instruction, multiple data - is _____ possible but programmers normally write a _____ that runs on all processors relying on _____ statements when _____ processors should execute _____ sections of code. This style is single _____, multiple data.
- SIMD – single instruction, multiple data – operate on _____ of data. SIMD needs only _____ of the code that is being simultaneously executed. SIMD works best when dealing with _____ in _____ loops.

- The _____ that inspired the SIMD category faded into history but two _____ interpretations of SIMD remain _____ today.
 - _____
 - _____

6.3 x86 Multimedia Extensions

- The most _____ used variation of SIMD is the basis of the hundreds of MMX and SSE instructions of the x86 processor. These instructions were added to improve performance of _____ programs.
 - The hardware allows flexible ALU operations – one 64-bit or two 32-bit or for 16-bit or eight 8-bit
 - Loads and stores are simply as wide as the widest ALU.
 - SSE now supports simultaneous execution of a pair of 64-bit floating-point numbers

6.3 Vector

- An older and more elegant interpretation of SIMD is called a vector architecture, which has been closely identified with Cray Computers.
- Consider $Y = a \times X + Y$

Original

```
l.d    $f0, a($sp)
addiu  $t1, $s0, #512
loop:  l.d    $f2, 0($s0)
mul.d  $f2, $f2, $f0
l.d    $f4, 0($s1)
add.d  $f4, $f4, $f2
s.d    $f4, 0($s1)
addiu  $s0, $s1, #8
addiu  $s1, $s1, #8
subu   $t0, $t1, $s0
bne    $t0, $zero, loop
```

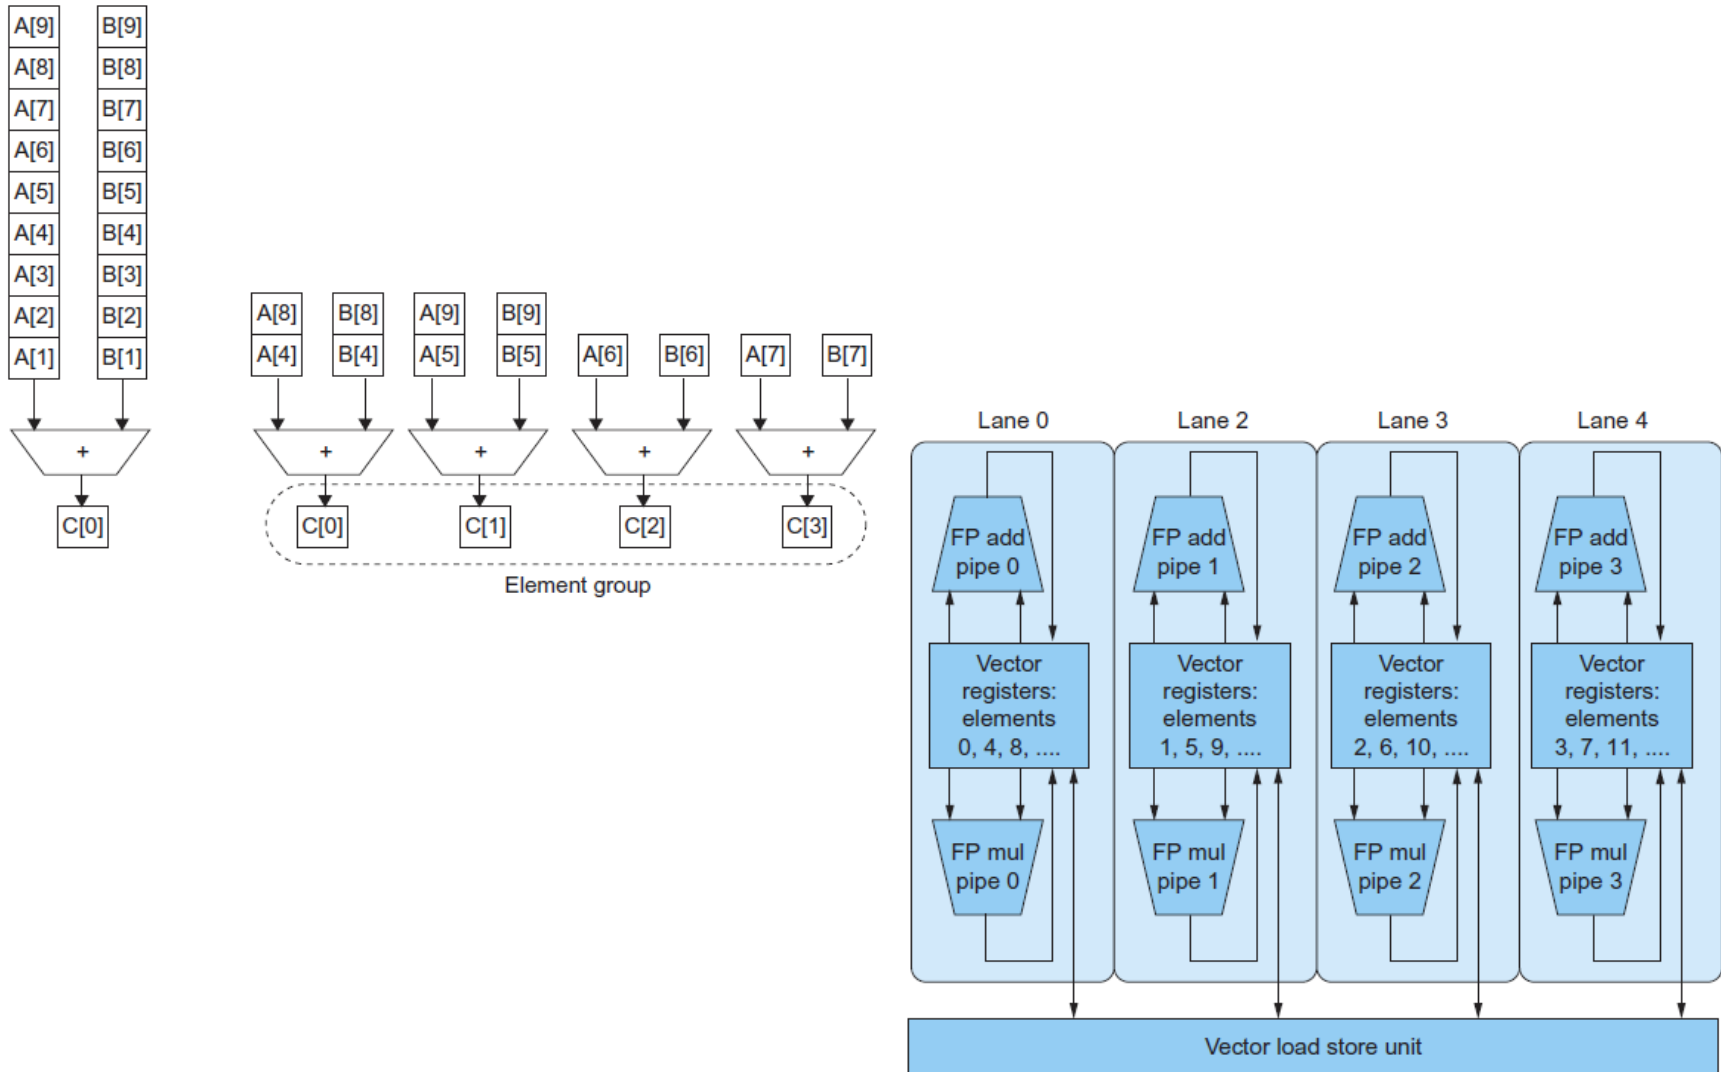
Vector

```
l.d    $f0, a($sp)
lv      $v1, 0($s0)
mulvs.d $v2, $v1, $f0
lv      $v3, 0($s1)
addv.d  $v4, $v4, $v3
sv      $v4, 0($s1)
```

6.3 Comparisons

- Vector versus Scalar
 - A _____ instruction specifies a great deal of work, the instruction _____ and _____ bandwidth is greatly reduced.
 - Hardware does not have to check for _____ a vector instruction.
 - Vector architectures and compilers have worked well for _____.
 - Hardware need only check for _____ hazards _____ between two vector _____, not _____ for every _____.
 - Vector instructions that access memory have a _____, memory system can be adjusted accordingly.
 - Replacing a _____ with a _____ reduces _____ hazards.
- Vector versus Multimedia Extensions
 - Vector specifies the number of operands in _____, not in _____.
 - Vector data transfers need _____ be _____.
 - Vector _____ over time more _____.

6.3 Improving the Performance of Vector




6.4 Hardware Multithreading

- Hardware multithreading allows _____ to _____ the _____ units of a single processor in _____ fashion.
- The processor must _____ the independent _____ of each _____.
- The hardware must support the _____ to _____ to different _____ relatively quickly.
- _____ multithreading switches between _____ on each _____, often done round robin.
 - Hides _____ losses by doing useful work during _____.
 - Inserts _____ for threads with _____.
- _____ multithreading _____ threads only on _____ stalls, such as _____ misses.
 - It is limited in its ability to overcome _____ losses, especially from _____ stalls. The major problem is pipeline _____.

6.4 Simultaneous Multithreading

- Simultaneous multithreading uses the resources of a _____, _____ scheduled processor to exploit _____ parallelism at the same time it exploits _____ parallelism.
- The key insight is that multiple-issue processors often have more _____ than a _____ thread can effectively use.
- _____ of dependences can be handled by the _____ capability.



6.4 Multithreading Speedup





FIGURE 6.6 The speed-up from using multithreading on one core on an i7 processor averages 1.31 for the PARSEC benchmarks (see [Section 6.9](#)) and the energy efficiency improvement is 1.07. This data was collected and analyzed by Esmailzadeh et. al. [2011].

6.5 Multicore and Other Shared Memory Multiprocessors

- A shared memory multiprocessor (SMP) is one that offers the programmer a _____ across all processors
- Processor communicate through _____ in memory.
- SMPs come in two flavors
 - _____
 - _____
- Processors need to coordinate when sharing data, this process is called _____, processors must acquire a _____




6.5 Shared Address Space Parallel Program (1)

- Suppose we want to sum 64,000 numbers on an SMP with UMA. Let's assume we have 64 processors.

```

sum[Pn] = 0;
for (i = 1000*Pn; i < 1000*(Pn+1); i = i + 1)
    sum[Pn] = sum[Pn] + A[i]; /* sum the assigned areas */
  
```

- After execution of this code, there are 64 _____ sums
- Need to _____ them into _____ sum
- Do so using a _____



6.5 Shared Address Space Parallel Program (2)

```
half = 64; /* 64 processors in multiprocessor */
repeat
    synch(); /* wait for partial sum completion */
    if (half%2 != 0 && Pn == 0)
        sum[0] = sum[0] + sum[half-1];
        /* Conditional sum needed when half is odd;
           Processor0 gets missing element */
    half = half/2; /* dividing line on who sums */
    if (Pn < half) sum[Pn] = sum[Pn] + sum[Pn+ half];
until (half == 1); /* exit with final sum in sum[0] */
```

6.5 A Parallel Programming System

- A _____ but _____ example is _____
 - _____ is an _____ _____
along with a set of _____,
_____ variables, and _____ library routines.
 - It offers a _____, _____, and _____
programming model for shared memory multiprocessors.
 - Its primary goal is to _____ and to
perform _____.
 - _____ extends _____ using _____, commands to the C
macro processor

6.5 OpenMP Example

```
Cc -fopenmp foo.c
#define P 64 /* define a constant */
#pragma omp parallel num_threads(P)

#pragma omp parallel for
for (Pn = 0; Pn < P; Pn +=1)
    for (1000*Pn; i < 1000*(Pn +1); i +=1)
        sum[Pn] += A[i]; /* sum the assigned areas */

#pragma omp parallel for reduction(+: FinalSum)
for (i = 0; i < P; i += 1)
    FinalSum += sum[i]; /* Reduce to a single number */
```

6.6 Introduction to Graphics Processing Units

- A major driving force for improving graphics processing was the _____, a different _____ than the one for CPUs.
- Key differences between GPUs and CPUs
 - GPUs are _____ that _____ a CPU, they don't have to do _____.
 - GPU problems sizes are typically hundreds of _____ to _____, but not hundreds of _____ to _____.
- Different Architecture Features
 - GPUs do not rely on _____ caches, they rely on having enough _____ to _____ memory latency.
 - The GPU main memory is oriented towards _____ rather than _____.
 - Each GPU processor is more highly _____ than a typical CPU, plus they have _____.

6.6 Programming GPUs

- Initially programmers had only _____ and _____.
 - They developed _____ languages
 - NVIDIA _____ (CUDA)
 - OpenCL is a _____ initiative to develop a _____ programming language
 - Unifying theme is _____
 - _____ and _____ can gang thousands of CUDA threads together to utilize _____, _____, _____, and _____
 - Threads are _____ together and executed in _____ of _____ at a time
 - A _____ processor inside a GPU executes these blocks of threads, and a GPU consists of _____ to _____ of these _____ processors.