

Routing

Chapter 3

Forwarding versus Routing

- Forwarding:
 - to select an output port based on destination address and routing table
- Routing:
 - process by which routing table is built



1

Routing

Chapter 3

- Forwarding table VS Routing table
 - Forwarding table
 - Used when a packet is being forwarded and so must contain enough information to accomplish the forwarding function
 - A row in the forwarding table contains the mapping from a network number to an outgoing interface and some MAC information, such as Ethernet Address of the next hop
 - Routing table
 - Built by the routing algorithm as a precursor to build the forwarding table
 - Generally contains mapping from network numbers to next hops



2

Routing

Chapter 3

(a)		
Prefix/Length	Next Hop	
18/8	171.69.245.10	


(b)		
Prefix/Length	Interface	MAC Address
18/8	if0	8:0:2b:e4:b:1:2

Example rows from (a) routing and (b) forwarding tables

Routing

Chapter 3

- Network as a Graph



- The basic problem of routing is to find the lowest-cost path between any two nodes
- The cost of a path equals the sum of the costs of all the edges that make up the path

Routing

Chapter 3

- For a simple network, we can calculate all shortest paths and load them into some nonvolatile storage on each node.
- Such a static approach has several shortcomings
 - It does not deal with node or link failures
 - It does not consider the addition of new nodes or links
 - It implies that edge costs cannot change
- What is the solution?
 - Need a distributed and dynamic protocol
 - Two main classes of protocols
 - Distance Vector
 - Link State


MK
MORGAN KAUFMANN

5

Distance Vector Routing Algorithm

Chapter 3

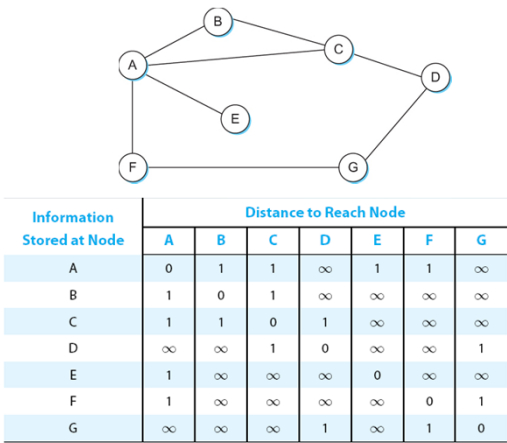
- Each node constructs a one dimensional array (a vector) containing the “distances” (costs) to all other nodes and distributes that vector to its immediate neighbors
- Starting assumption is that each node knows the cost of the link to each of its directly connected neighbors

MK
MORGAN KAUFMANN

6

Distance Vector - Example

Chapter 3



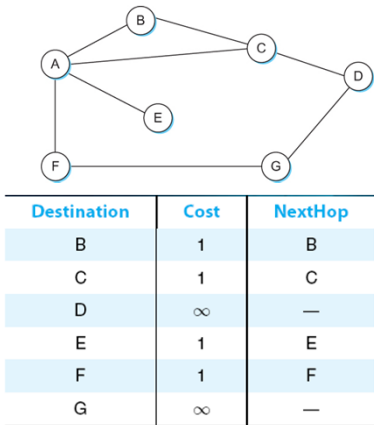
Initial distances stored at each node (global view)
Use hop count as the cost



7

Distance Vector - Example

Chapter 3

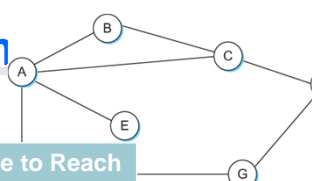


Initial routing table at node A



8

Distance Vector - Exam



Information Stored at Node	Distance to Reach Node						
	A	B	C	D	E	F	G
A	0	1	1	∞	1	1	∞
B	1	0	1	∞	∞	∞	∞
C	1	1	0	1	∞	∞	∞
D	∞	∞	1	0	∞	∞	1
E	1	∞	∞	∞	0	∞	∞
F	1	∞	∞	∞	∞	0	1
G	∞	∞	∞	1	∞	1	0

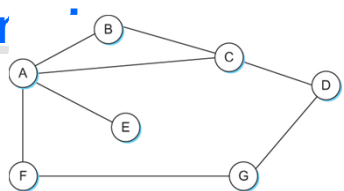
Information Stored at Node	Distance to Reach Node						
	A	B	C	D	E	F	G
A	0	1	1	2	1	1	2
B	1	0	1	2	2	2	∞
C	1	1	0	1	2	2	2
D	2	2	1	0	∞	2	1
E	1	2	2	∞	0	2	∞
F	1	2	2	2	2	0	1
G	2	∞	2	1	∞	1	0

Distances stored at each node after 1 update (global view)

MK

9

Distance Vector - Exam



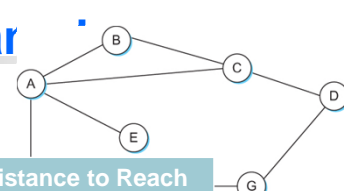
Information Stored at Node	Distance to Reach Node						
	A	B	C	D	E	F	G
A	0	1	1	2	1	1	2
B	1	0	1	2	2	2	∞
C	1	1	0	1	2	2	2
D	2	2	1	0	∞	2	1
E	1	2	2	∞	0	2	∞
F	1	2	2	2	2	0	1
G	2	∞	2	1	∞	1	0

Distances stored at each node after 1 update (global view)

MK

10

Distance Vector - Exam



Information Stored at Node	Distance to Reach Node						
	A	B	C	D	E	F	G
A	0	1	1	2	1	1	2
B	1	0	1	2	2	2	∞
C	1	1	0	1	2	2	2
D	2	2	1	0	∞	2	1
E	1	2	2	∞	0	2	∞
F	1	2	2	2	2	0	1
G	2	∞	2	1	∞	1	0

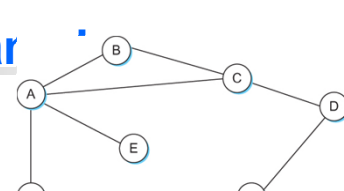
Information Stored at Node	Distance to Reach Node						
	A	B	C	D	E	F	G
A	0	1	1	2	1	1	2
B	1	0	1	2	2	2	3
C	1	1	0	1	2	2	2
D	2	2	1	0	3	2	1
E	1	2	2	3	0	2	3
F	1	2	2	2	2	0	1
G	2	3	2	1	3	1	0

Distances stored at each node after 2 updates (global view)

MK
MORGAN KAUFMANN

11

Distance Vector - Exam



Information Stored at Node	Distance to Reach Node						
	A	B	C	D	E	F	G
A	0	1	1	2	1	1	2
B	1	0	1	2	2	2	3
C	1	1	0	1	2	2	2
D	2	2	1	0	3	2	1
E	1	2	2	3	0	2	3
F	1	2	2	2	2	0	1
G	2	3	2	1	3	1	0

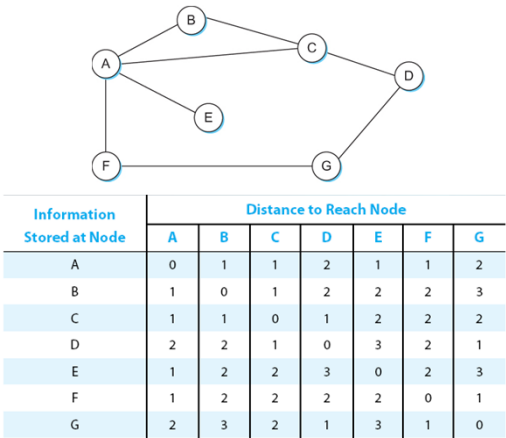
Distances stored at each node after 2 updates (global view)

MK
MORGAN KAUFMANN

12

Distance Vector - Example

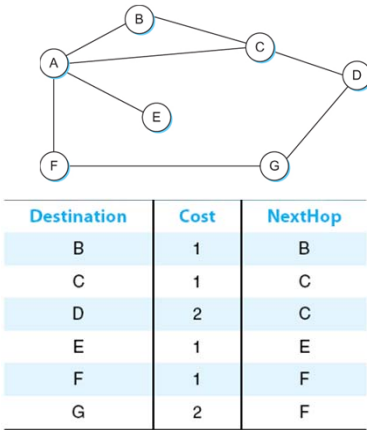
Chapter 3



Final distances stored at each node (global view)

Distance Vector - Example

Chapter 3




Final routing table at node A

Distance Vector Routing Algorithm

- The distance vector routing algorithm is sometimes called as Bellman-Ford algorithm
- Every T seconds each router sends its table to its neighbor each router then updates its table based on the new information
- Problems include fast response to good news and slow response to bad news. Also too many messages to update


Distance Vector Routing Algorithm

- When a node detects a link failure
 - F detects that link to G has failed
 - F sets distance to G to infinity and sends update to A
 - A sets distance to G to infinity since it uses F to reach G
 - A receives periodic update from C with 2-hop path to G
 - A sets distance to G to 3 and sends update to F
 - F decides it can reach G in 4 hops via A




Distance Vector Routing Algorithm

- Slightly different circumstances can prevent the network from stabilizing
 - Suppose the link from A to E goes down
 - In the next round of updates, A advertises a distance of infinity to E, but B and C advertise a distance of 2 to E



Distance Vector Routing Algorithm

- Depending on the exact timing of events, the following might happen
 - Node B, upon hearing that E can be reached in 2 hops from C, concludes that it can reach E in 3 hops and advertises this to A
 - Node A concludes that it can reach E in 4 hops and advertises this to C
 - Node C concludes that it can reach E in 5 hops; and so on.
 - This cycle stops only when the distances reach some number that is large enough to be considered infinite
 - Count-to-infinity problem**



Chapter 3

Count-to-infinity Problem

- Use some relatively small number as an approximation of infinity – i.e 16
- Assign a maximum number of hops to get across a certain network

MK
MORGAN KAUFMANN

19

Chapter 3

Count-to-infinity Problem

- One technique to improve the time to stabilize routing is called *split horizon*
 - When a node sends a routing update to its neighbors, it does not send those routes it learned from each neighbor back to that neighbor
 - For example, if B has the route (E, 2, A) in its table, then it knows it must have learned this route from A, and so whenever B sends a routing update to A, it does not include the route (E, 2) in that update


MK
MORGAN KAUFMANN

20

Count-to-infinity Problem

Chapter 3


- In a stronger version of split horizon, called *split horizon with poison reverse*
 - B actually sends that route back to A, but it puts negative information in the route to ensure that A will not eventually use B to get to E
 - For example, B sends the route (E, ∞) to A



21

Routing Information Protocol (RIP)

Chapter 3



Example Network running RIP

8		16		31	
Command	Version		Must be zero		
Family of net 1			Route Tags		
Address prefix of net 1					
Mask of net 1					
Distance to net 1					
Family of net 2			Route Tags		
Address prefix of net 2					
Mask of net 2					
Distance to net 2					

RIPv2 Packet Format

Built on Distance Vector Protocol
Simple to implement – Routers advertise cost (# hops) to a network
Infinity set to 16 – limited to small internetworks



22

Link State Routing

Chapter 3

Strategy: Send to all nodes (not just neighbors) information about directly connected links (not entire routing table).

- Link State Packet (LSP)
 - id of the node that created the LSP
 - cost of link to each directly connected neighbor
 - sequence number (SEQNO)
 - time-to-live (TTL) for this packet
- Reliable Flooding
 - store most recent LSP from each node
 - forward LSP to all nodes but one that sent it
 - generate new LSP periodically; increment SEQNO
 - start SEQNO at 0 when reboot
 - decrement TTL of each stored LSP; discard when TTL=0




23

Link State

Chapter 3

Reliable Flooding



Flooding of link-state packets. (a) LSP arrives at node X; (b) X floods LSP to A and C; (c) A and C flood LSP to B (but not X); (d) flooding is complete



24

Shortest Path Routing

- Dijkstra's Algorithm - Assume non-negative link weights
 - N : set of nodes in the graph
 - $l(i, j)$: the non-negative cost associated with the edge between nodes $i, j \in N$ and $l(i, j) = \infty$ if no edge connects i and j
 - Let $s \in N$ be the starting node which executes the algorithm to find shortest paths to all other nodes in N
 - Two variables used by the algorithm
 - M : set of nodes incorporated so far by the algorithm
 - $C(n)$: the cost of the path from s to each node n
 - The algorithm


```

M = {s}
For each n in N - {s}
    C(n) = l(s, n)
while ( N ≠ M)
    M = M ∪ {w} such that C(w) is the minimum
                        for all w in (N-M)

    For each n in (N-M)
        C(n) = MIN (C(n), C(w) + l(w, n))
              
```

Shortest Path Routing

- In practice, each switch computes its routing table directly from the LSP's it has collected using a realization of Dijkstra's algorithm called the *forward search algorithm*
- Specifically each switch maintains two lists, known as **Tentative** and **Confirmed**
- Each of these lists contains a set of entries of the form (Destination, Cost, NextHop)


Shortest Path Routing-Forward Search

Chapter 3

- The algorithm
 - 1) Initialize the **Confirmed** list with an entry for myself; this entry has a cost of 0
 - 2) For the node just added to the **Confirmed** list in the previous step, call it node **Next**, select its LSP
 - 3) For each neighbor (Neighbor) of **Next**, calculate the cost (Cost) to reach this Neighbor as the sum of the cost from myself to Next and from Next to Neighbor
 - a) If Neighbor is currently on neither the **Confirmed** nor the **Tentative** list, then add (Neighbor, Cost, Nexthop) to the **Tentative** list, where Nexthop is the direction I go to reach Next
 - b) If Neighbor is currently on the **Tentative** list, and the Cost is less than the currently listed cost for the Neighbor, then replace the current entry with (Neighbor, Cost, Nexthop) where Nexthop is the direction I go to reach Next
 - c) If Neighbor on **Confirmed** list, then ignore it
 - 4) If the **Tentative** list is empty, stop. Otherwise, pick the entry from the **Tentative** list with the lowest cost, move it to the **Confirmed** list, and return to Step 2.

Shortest Path Routing-Forward Search

Chapter 3



D has LSP's of all nodes (Destination, Cost, NextHop)

Step	Confirmed	Tentative	Comments
1	(D,0,-)		Since D is the only new member of the confirmed list, look at its LSP.
2	(D,0,-)	(B,11,B) (C,2,C)	D's LSP says we can reach B through B at cost 11, which is better than anything else on either list, so put it on Tentative list; same for C.
3	(D,0,-) (C,2,C)	(B,11,B)	Put lowest-cost member of Tentative (C) onto Confirmed list. Next, examine LSP of newly confirmed member (C).
4	(D,0,-) (C,2,C)	(B,5,C) (A,12,C)	Cost to reach B through C is 5, so replace (B,11,B). C's LSP tells us that we can reach A at cost 12.
5	(D,0,-) (C,2,C) (B,5,C)	(A,12,C)	Move lowest-cost member of Tentative (B) to Confirmed, then look at its LSP.
6	(D,0,-) (C,2,C) (B,5,C)	(A,10,C)	Since we can reach A at cost 5 through B, replace the Tentative entry.
7	(D,0,-) (C,2,C) (B,5,C) (A,10,C)		Move lowest-cost member of Tentative (A) to Confirmed, and we are all done.

Open Shortest Path First (OSPF)

Chapter 3

0	8	16	31
Version	Type	Message length	
SourceAddr			
Areald			
Checksum		Authentication type	
Authentication			

OSPF Header Format

LS Age		Options	Type = 1
Link-state ID			
Advertising router			
LS sequence number			
LS checksum		Length	
0	Flags	0	Number of links
Link ID			
Link data			
Link type	Num_TOS	Metric	
Optional TOS information			
More links			

OSPF Link State Advertisement

Distance-Vector vs Link-State

Chapter 3

- **Distance-Vector(RIP):**
 - Each node talks with its neighbors only
 - Sends all information it knows - known distances to other nodes
 - Speed of convergence is slower than LS
 - Stabilization may not occur – count to infinity
 - Simple algorithm
- **Link-State (OSPF):**
 - Each node talks to all other nodes
 - Sends what it knows for sure - State of its directly connected links
 - Stabilizes quickly and it responds rapidly to network changes
 - Low traffic generation
 - Storage at each node is large
 - Uses reliable flooding of packets

Metrics – Cost of Links

- Assign 1 to all links – hop count
 - Latency – take into account delay of the link
 - Capacity – what is BW of each link
 - Current Load – increase cost as load increases
 - Queue length (average value between updates)
-
- Metrics are fixed by administrators – not dynamically changing due to stability issues.

Summary

- We have looked at some of the issues involved in building scalable and heterogeneous networks by using switches and routers to interconnect links and networks.
- To deal with heterogeneous networks, we have discussed in details the service model of Internetworking Protocol (IP) which forms the basis of today's routers.
- We have discussed in details two major classes of routing algorithms
 - Distance Vector
 - Link State