

Application Programming Interface (Sockets)

- Socket Interface was originally provided by the Berkeley distribution of Unix
 - Now supported in virtually all operating systems
- Each protocol provides a certain set of *services*, and the API provides a syntax by which those services can be invoked in this particular OS

Sockets

- What is a socket?
 - The point where a local application process attaches to the network
 - An interface between an application and the network
 - An application creates the socket
- The interface defines operations for
 - Creating a socket
 - Attaching a socket to the network
 - Sending and receiving messages through the socket
 - Closing the socket

Sockets

Chapter 1

- Socket Family
 - PF_INET denotes the Internet family
 - PF_UNIX denotes the Unix pipe facility
 - PF_PACKET denotes direct access to the network interface (i.e., it bypasses the TCP/IP protocol stack)
- Socket Type
 - SOCK_STREAM is used to denote a byte stream
 - SOCK_DGRAM is an alternative that denotes a message oriented service, such as that provided by UDP



3

Sockets

Chapter 1

- Socket Protocol – specifies protocol to use
 - Non zero value specifies a protocol supported by the specified socket family
 - 0 invokes the default protocol used by the socket family and type specified.



4

Creating a Socket

```
int sockfd = socket(address_family, type, protocol);
```

- The socket number returned is the socket descriptor (sockfd) for the newly created socket
- `int sockfd = socket (PF_INET, SOCK_STREAM, 0);`
- `int sockfd = socket (PF_INET, SOCK_DGRAM, 0);`

The combination of PF_INET and SOCK_STREAM implies TCP

The combination of PF_INET and SOCK_DGRAM implies UDP

Client-Serve Model with TCP

Server

- Passive open
- Prepares to accept connection, does not actually establish a connection

Server invokes

```
int bind (int socket, struct sockaddr *address,
          int addr_len)

int listen (int socket, int backlog)

int accept (int socket, struct sockaddr *address,
            int *addr_len)
```

Client-Serve Model with TCP: Bind - Server

Chapter 1

```
int bind (int socket, struct sockaddr *address, int addr_len)
```

- Binds the newly created socket to the specified address i.e. the network address of the local participant (the server)
- **int socket**: file descriptor of the socket being bound (from socket call)
- **struct sockaddr *address** is a data structure which combines IP and port
- **int addr_len** specifies the length of the sockaddr structure
- Upon successful completion, returns 0. Unsuccessful completion returns -1



7

Client-Serve Model with TCP: Listen-Server

Chapter 1

```
int listen (int socket, int backlog)
```

- **int socket**: file descriptor of the socket being bound (from socket call)
- **int backlog**: number of outstanding connections in the sockets listen queue. Defines how many connections can be pending on the specified socket
- Return value of 0 indicates success, return value of -1 indicates an error occurred



8

Client-Serve Model with TCP: Accept-Server

Chapter 1

- Accept
 - Carries out the passive open
 - Blocking operation
 - Does not return until a remote participant has established a connection
 - When it does, it returns a new socket that corresponds to the new established connection and the address argument contains the remote participant's address



9

Client-Serve Model with TCP: Accept-Server

Chapter 1

```
int accept (int socket, struct sockaddr
            *address, int *addr_len)
```

- **int socket**: file descriptor of the socket being bound (from socket call)
- **struct sockaddr *address** is a data structure which combines IP and port
- **int *addr_len** sends in the length of the supplied address, on output from the function it contains the length of the stored address
- Return value on successful acceptance is the file descriptor of the accepted socket. Return value of -1 indicates an error occurred.



10

Client-Serve Model with TCP

Client

- Application performs active open
- It says who it wants to communicate with

Client invokes

```
int connect (int socket, struct sockaddr *address,
            int addr_len)
```

Connect

- Does not return until TCP has successfully established a connection at which time application is free to begin sending data
- Address contains remote machine's address

Client-Serve Model with TCP: Connect-client

```
int connect (int socket, struct sockaddr *address,
            int addr_len)
```

- **int socket**: file descriptor of the socket being bound (from socket call)
- **struct sockaddr *address** is a data structure which combines IP and port – address of machine to connect to
- **int *addr_len** the length of the supplied address
- Return value of 0 on successful completion. Return value of -1 indicates an error occurred.

Client-Serve Model with TCP

In practice

- The client usually specifies only remote participant's address and let's the system fill in the local information
- Whereas a server usually listens for messages on a well-known port
- A client does not care which port it uses for itself, the OS simply selects an unused one

Client-Serve Model with TCP

Once a connection is established, the application process invokes two operation

```
int send (int socket, char *msg, int msg_len,  
          int flags)
```

```
int recv (int socket, char *buff, int buff_len,  
          int flags)
```

Client-Serve Model with TCP - send

```
int send (int socket, char *msg, int msg_len,  
          int flags)
```

- **int socket**: file descriptor of the socket being bound (from socket call)
- **char *msg** points to the buffer containing the message to send
- **int msg_len** the length of the message in bytes
- **Int flags** specifies type of message transmission
- Return value of -1 if a local error is detected. No guarantee of delivery status.
- If message is too long to send via the specified protocol, send fails and no information is sent

Client-Serve Model with TCP - recv

```
int recv (int socket, char *buff, int buff_len,  
          int flags)
```

- **int socket**: file descriptor of the socket being bound (from socket call)
- **char *buff** points to the buffer where the message received is to be stored
- **int buff_len** the length of the storage buffer in bytes
- **Int flags** specifies type of message reception
- Return value is the length of the message in bytes. Return value of 0 if no messages are available and a shutdown has been received. Return value of -1 if an error is detected. No guarantee of delivery status.

Example Application: Client

```
#include <stdio.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <netdb.h>

#define SERVER_PORT 5432
#define MAX_LINE 256

int main(int argc, char * argv[])
{
    FILE *fp;
    struct hostent *hp;
    struct sockaddr_in sin;
    char *host;
    char buf[MAX_LINE];
    int s;
    int len;
    if (argc==2) {
        host = argv[1];
    }
    else {
        fprintf(stderr, "usage: simplex-talk host\n");
        exit(1);
    }
}
```

Example Application: Client

```
/* translate host name into peer's IP address */
hp = gethostbyname(host);
if (!hp) {
    fprintf(stderr, "simplex-talk: unknown host: %s\n", host);
    exit(1);
}
/* build address data structure */
bzero((char *)&sin, sizeof(sin));
sin.sin_family = AF_INET;
bcopy(hp->h_addr, (char *)&sin.sin_addr, hp->h_length);
sin.sin_port = htons(SERVER_PORT);
/* active open */
if ((s = socket(PF_INET, SOCK_STREAM, 0)) < 0) {
    perror("simplex-talk: socket");
    exit(1);
}
if (connect(s, (struct sockaddr *)&sin, sizeof(sin)) < 0) {
    perror("simplex-talk: connect");
    close(s);
    exit(1);
}
/* main loop: get and send lines of text */
while (fgets(buf, sizeof(buf), stdin)) {
    buf[MAX_LINE-1] = '\0';
    len = strlen(buf) + 1;
    send(s, buf, len, 0);
}
}
```

Example Application: Server

```
#include <stdio.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <netdb.h>
#define SERVER_PORT 5432
#define MAX_PENDING 5
#define MAX_LINE 256

int main()
{
    struct sockaddr_in sin;
    char buf[MAX_LINE];
    int len;
    int s, new_s;
    /* build address data structure */
    bzero((char *)&sin, sizeof(sin));
    sin.sin_family = AF_INET;
    sin.sin_addr.s_addr = INADDR_ANY;
    sin.sin_port = htons(SERVER_PORT);

    /* setup passive open */
    if ((s = socket(PF_INET, SOCK_STREAM, 0)) < 0) {
        perror("simplex-talk: socket");
        exit(1);
    }
}
```

Example Application: Server

```
if ((bind(s, (struct sockaddr *)&sin, sizeof(sin))) < 0) {
    perror("simplex-talk: bind");
    exit(1);
}
listen(s, MAX_PENDING);
/* wait for connection, then receive and print text */
while(1) {
    if ((new_s = accept(s, (struct sockaddr *)&sin, &len)) < 0) {
        perror("simplex-talk: accept");
        exit(1);
    }
    while (len = recv(new_s, buf, sizeof(buf), 0))
        fputs(buf, stdout);
    close(new_s);
}
}
```

Socket Web pages

- The following are web pages for the various aspects of socket function descriptions:
 - <http://pubs.opengroup.org/onlinepubs/009695399/functions/socket.html>
 - <http://pubs.opengroup.org/onlinepubs/009695399/functions/bind.html>
 - <http://pubs.opengroup.org/onlinepubs/009695399/functions/send.html>
 - <http://pubs.opengroup.org/onlinepubs/009695399/functions/listen.html>
 - <http://pubs.opengroup.org/onlinepubs/009695399/functions/accept.html>
 - <http://pubs.opengroup.org/onlinepubs/009695399/functions/recv.html>
 - <http://pubs.opengroup.org/onlinepubs/009695399/functions/connect.html>
 - <http://beej.us/guide/bgnet/output/html/multipage/htonsman.html>
 - Following is a complete tutorial
 - <http://beej.us/guide/bgnet/output/html/multipage/index.html>