

Reliable Transmission

- CRC is used to detect errors.
- Some error codes are strong enough to correct errors.
- The overhead is typically too high.
- Corrupt frames must be discarded.
- A link-level protocol that wants to deliver frames reliably must recover from these discarded frames.
- This is accomplished using a combination of two fundamental mechanisms
 - Acknowledgements and Timeouts

Reliable Transmission

- An *acknowledgement* (ACK for short) is a small control frame that a protocol sends back to its peer saying that it has received the earlier frame.
 - A control frame is a frame with header only (no data).
- The receipt of an *acknowledgement* indicates to the sender of the original frame that its frame was successfully delivered.


Reliable Transmission

- If the sender does not receive an *acknowledgment* after a reasonable amount of time, then it retransmits the original frame.
- The action of waiting a reasonable amount of time is called a *timeout*.
- The general strategy of using *acknowledgements* and *timeouts* to implement reliable delivery is sometimes called **Automatic Repeat reQuest (ARQ)**.

Stop and Wait Protocol

- Idea of stop-and-wait protocol is straightforward
 - After transmitting one frame, the sender waits for an acknowledgement before transmitting the next frame.
 - If the acknowledgement does not arrive after a certain period of time, the sender times out and retransmits the original frame

Stop and Wait Protocol



Timeline showing four different scenarios for the stop-and-wait algorithm.


(a) The ACK is received before the timer expires; (b) the original frame is lost; (c) the ACK is lost; (d) the timeout fires too soon

Stop and Wait Protocol

- If the acknowledgment is lost or delayed in arriving
 - The sender times out and retransmits the original frame,
 - Receiver will think that it is the next frame since it has correctly received and acknowledged the first frame
 - As a result, duplicate copies of frames will be delivered
- How to solve lost ACK problem
 - Use 1 bit sequence number (0 or 1)
 - The sender retransmits frame 0,
 - The receiver can determine that it is seeing a second copy of frame 0 rather than the first copy of frame 1 and therefore can ignore it
 - the receiver still acknowledges it, in case the first ACK was lost

Stop and Wait Protocol

Chapter 2



Timeline for stop-and-wait with 1-bit sequence number

MK
MORGAN KAUFMANN

7

Stop and Wait Protocol - issue

Chapter 2

- With Stop and Wait, The sender has only one outstanding frame on the link at a time
 - This may be far below the link's capacity – poor utilization
 - Sending rate = (bits per frame)/(time per frame = 1 RTT)


MK
MORGAN KAUFMANN

8

Stop and Wait Protocol - issue

- Consider sending a 1 KB (1024 Bytes) frame over a 1.5 Mbps link with a 45 ms RTT
 - The link has a delay \times bandwidth product of 67.5 Kb or approximately 8 KB ($1.5\text{E}6 \times 0.045 = 67,500 \text{ bits} = 8,437 \text{ Bytes}$)
 - Since the sender can send only one frame per RTT
 - Maximum Sending rate is
 $\text{Bits per frame} \div \text{Time per frame} = 1024 \times 8 \div 0.045 = 182 \text{ Kbps}$
 Or about one-eighth of the link's capacity
 - To use the link fully, then sender should transmit up to eight frames before having to wait for an acknowledgement

Sliding Window Protocol



Sliding Window Protocol - Sender

Chapter 2

- Sender assigns a sequence number denoted as SeqNum to each frame.
 - Assume it can grow infinitely large
- Sender maintains three variables
 - Sending Window Size (SWS)
 - Upper bound on the number of outstanding (unacknowledged) frames that the sender can transmit
 - Last Acknowledgement Received (LAR)
 - Sequence number of the last acknowledgement received
 - Last Frame Sent (LFS)
 - Sequence number of the last frame sent




11

Sliding Window Protocol - Sender

Chapter 2

- Sender also maintains the following invariant (property)

$$\text{LFS} - \text{LAR} \leq \text{SWS}$$



Sliding Window on Sender

- LAR moves right as ACKs received
- LFS moves right as frames are sent



12

Sliding Window Protocol - Sender

Chapter 2

- When an acknowledgement arrives
 - the sender moves LAR to right
 - Allows the sender to transmit another frame
- The sender associates a timer with each frame it transmits
 - It retransmits the frame if the timer expires before the ACK is received
- Note that the sender has to be willing to buffer up to SWS frames
 - WHY?



13

Sliding Window Protocol - Rcvr

Chapter 2

- Receiver maintains three variables
 - Receiving Window Size (RWS)
 - Upper bound on the number of out-of-order frames that the receiver is willing to accept
 - Largest Acceptable Frame (LAF)
 - Sequence number of the largest acceptable frame
 - Last Frame Received (LFR)
 - Sequence number of the last frame received and acknowledged




14

Sliding Window Protocol - Rcvr

- Receiver also maintains the following invariant

$$\text{LAF} - \text{LFR} \leq \text{RWS}$$



Sliding Window on Receiver

Sliding Window Protocol - Rcvr

- When a frame with sequence number SeqNum arrives, what does the receiver do?
 - If $\text{SeqNum} \leq \text{LFR}$ or $\text{SeqNum} > \text{LAF}$
 - Discard it (the frame is outside the receiver window)
 - If $\text{LFR} < \text{SeqNum} \leq \text{LAF}$
 - Accept it
 - Now the receiver needs to decide whether or not to send an ACK

Sliding Window Protocol - Rcvr

- Let SeqNumToAck
 - Denote the largest sequence number *not yet acknowledged*,
 - All frames with sequence number less than SeqNumToAck have been received
 - When Frame with sequence number SeqNumToAck is received, an ACK is sent for that frame
- The receiver acknowledges the receipt of SeqNumToAck even if higher-numbered packets have been received
 - This acknowledgement is said to be cumulative.
- The receiver then sets
 - $LFR = SeqNumToAck$ and adjusts
 - $LAF = LFR + RWS$

Sliding Window Protocol

For example, suppose $LFR = 5$ and $RWS = 4$
 (i.e. the last ACK that the receiver sent was for seq. no. 5 and SeqNumToAck is set to 6)
 → $LAF = 9$

If frames 7 and 8 arrive before frame 6, they will be buffered because they are within the receiver window
 But no ACK will be sent since frame 6 is yet to arrive
 Frames 7 and 8 are out of order

Frame 6 arrives (it is late because it was lost first time and had to be retransmitted)
 Receiver Acknowledges Frame 8
 Receiver bumps LFR to 8
 Receiver moves LAF to 12 ($LAF = LFR + RWS$)
 Receiver sets SeqNumToAck to 9

Issues with Sliding Window Protocol

- When timeout occurs, the amount of data in transit decreases
 - Sender is unable to advance its window
- When the packet loss occurs, this scheme is no longer keeping the pipe full
 - The longer it takes to notice that a packet loss has occurred, the more severe the problem becomes
- How to improve this
 - Negative Acknowledgement (NAK)
 - Additional Acknowledgement
 - Selective Acknowledgement

Issues with Sliding Window Protocol

- Negative Acknowledgement (NAK)
 - Receiver sends NAK for frame 6 when frame 7 arrive (in the previous example)
 - Unnecessary since sender's timeout mechanism will be sufficient to catch the situation
 - Adds additional complexity
- Additional Acknowledgement
 - Receiver sends additional ACK for frame 5 when frame 7 arrives
 - Sender uses duplicate ACK as a clue for frame loss
- Selective Acknowledgement
 - Receiver will acknowledge exactly those frames it has received, rather than the highest number frames
 - Receiver will acknowledge frames 7 and 8
 - Sender knows frame 6 is lost
 - Sender can keep the pipe full (additional complexity)

Issues with Sliding Window Protocol

How to select the window size

- SWS is easy to compute
 - Use $\text{Delay} \times \text{Bandwidth}/(\text{frame size})$ – keeps the pipe full
- RWS can be anything
 - Two common setting
 - $\text{RWS} = 1$

No buffer for frames that arrive out of order
 - $\text{RWS} = \text{SWS}$

The receiver buffers frames that the sender transmits
 - It does not make any sense to keep $\text{RWS} > \text{SWS}$. WHY?
Cannot have more than SWS frames arrive out of order

Issues with Sliding Window Protocol

- Finite Sequence Number
 - Frame sequence number is specified in the header field
 - Finite size

3 bits: eight possible sequence number: 0, 1, 2, 3, 4, 5, 6, 7
 - It is necessary to wrap around – reuse sequence numbers

Issues with Sliding Window Protocol

- How to distinguish between different incarnations of the same sequence number?
 - Number of possible sequence number must be larger than the number of outstanding frames allowed
 - Stop and Wait:
 - One outstanding frame → 2 distinct sequence number (0 and 1)
 - Let MaxSeqNum be the number of available sequence numbers
 - $SWS + 1 \leq \text{MaxSeqNum}$
 - Is this sufficient?

Issues with Sliding Window Protocol

- $SWS + 1 \leq \text{MaxSeqNum}$
 - Is this sufficient?
 - Depends on RWS
 - If $RWS = 1$, then sufficient
 - If $RWS = SWS$, then not good enough
- For example, we have eight sequence numbers
 0, 1, 2, 3, 4, 5, 6, 7
 $RWS = SWS = 7$
 Sender sends 0, 1, ..., 6
 Receiver receives 0, 1, ..., 6
 Receiver acknowledges 0, 1, ..., 6
 ACK (0, 1, ..., 6) are lost
 Sender retransmits 0, 1, ..., 6
 Receiver is expecting 7, 0, ..., 5

Issues with Sliding Window Protocol

- To avoid confusion on which packet has been received,
 - If $RWS = SWS$ (remember makes no sense for $RWS > SWS$)

$$SWS < (MaxSeqNum + 1)/2 \text{ or } MaxSeqNum > 2*SWS - 1$$

$MaxSeqNum$ is the number of sequence numbers required.

For Stop and Wait ARQ, $SWS = 1$, so $MaxSeqNum > 1$
 - If $RWS < SWS$, then $MaxSeqNum$ may be less than $2*SWS - 1$

Roles of Sliding Window Protocol

- Serves three different roles
 - Reliable
 - Preserve the order
 - Each frame has a sequence number
 - The receiver makes sure that it does not pass a frame up to the next higher-level protocol until it has already passed up all frames with a smaller sequence number
 - Frame control (flow control)
 - Receiver is able to throttle the sender
 - Keeps the sender from transmitting more data than the receiver is able to process
 - Can send how many more frames it can accept
- Sample code on pages 111-117